

KISHEN SHARMA

Daly City, CA | (858) 405-7873 | ksharma06@pm.me | [linkedin.com/in/kishen-sharma/](https://www.linkedin.com/in/kishen-sharma/) | github.com/ksharma6/

PROFESSIONAL EXPERIENCE

ML Engineer (Applied AI) | Inbox0 (<https://github.com/ksharma6/Inbox0>)

May 2025 – Present

- **Designed end-to-end agentic pipeline** integrating Gmail and Slack Bolt: defined prompt contracts and tool schemas to bind model outputs to executable Gmail actions and Slack interactions.
- **Architected production LangGraph agents** with structured tool-calling, schema-validated function execution, state checkpointing, retry logic, and human-in-the-loop approval gating for deterministic replay and production debugging.
- **Building a pipeline-wide evaluation harness** covering draft quality, triage accuracy, retrieval relevance, cost, latency, and token usage; integrates LLM-as-judge with pairwise comparison (position-bias aware), baselines components before changes ship, and ingests Slack action events plus Gmail folder state as continuous production signal.
- **Designing a tiered semantic retrieval system** (Redis working-set cache + SQLite-vec persistent store) addressing token growth observed during evaluation of long quoted-reply threads; architecture supports cross-thread retrieval over full email history with quality deltas measured through the evaluation harness.

Open-source Developer | Sktime (<https://www.sktime.net>)

Mar 2024 – Oct 2024

- **Contributed to a 1M+ user/month open-source ML framework:** decoupled GridSearch from scikit-learn, reducing user-reported issues by 20% and improving reproducibility across regression and classification workflows.
- **Shipped an AutoML MultiplexerRegressor** that auto-selects among state-of-the-art time-series regressors with GridSearch hyperparameter tuning, cutting model-selection time by 30% and simplifying downstream integration.

AI Training Engineer (RLHF Specialist) | Self-Employed

Feb 2024 – Nov 2024

- **OpenAI (ChatGPT):** evaluated LLM reasoning, tool usage, and prompt adherence across multi-turn conversations; wrote Python unit tests to verify generated code compiled correctly and matched user intent; flagged and corrected responses that diverged from prior context.
- **Google (Gemini):** rewrote chain-of-thought traces to compress agentic turn counts, evaluated tool-calling decisions, and validated production schemas for Maps, Search, Calendar, and Purchasing via code verification.
- **Meta (Llama):** scored output quality and applied heavy penalties for hallucinations and factually incorrect responses; regenerated outputs at the point of failure and evaluated whether downstream responses improved, producing preference data that reinforced correct reasoning trajectories.

Data Scientist | Acorn Analytics Inc.

Aug 2020 – Mar 2021

- **Automated AWS ETL pipelines with Python and MySQL** for data encryption and feature extraction; delivered analysis to C-suite that generated \$140K in new business opportunities.
- **Ran 10K Monte Carlo trials in a Markov chain framework** to generate hiring ranges and risk bands, briefed B2B SaaS executives and standardized a repeatable template later reused for Q1 and Q2 headcount planning.

Staff Consultant | Celerity Consulting Group

Mar 2016 – May 2018

- **Directed a team of five analysts, reducing burn rate 18%;** built weekly Tableau reporting via MySQL queries on SharePoint-hosted data (ETL, data modeling, scheduled refresh).
- **Collaborated with the client's Data Science team** to implement a logistic regression fraud-detection model resulting in \$250K cost savings; engineered domain-informed features that increased model lift and recall.

SKILLS

AI / LLM Systems: LangGraph, LangSmith, RAG, RLHF, Agent Orchestration, Prompt Engineering, Tool-Calling, LLM-as-Judge, Pairwise A/B Evaluation, Vector Retrieval (RAGAS, pgvector, SQLite-vec)

ML / Deep Learning: PyTorch, scikit-learn, Hugging Face Transformers

Languages: Python, SQL

Backend and Infrastructure: Redis, SQLite-vec, Flask, REST APIs, AWS, GCP, Linux, Unit/Integration Testing, CI/CD

EDUCATION

Stanford University

Jan 2022 – Jan 2024

- Graduate Coursework: CS229 (ML), CS224N (NLP), CS231N (Computer Vision), CS336 (LLMs), CS236 (GANs)

University of California, San Diego

June 2015

- Bachelor of Arts, Cognitive Science